

# iNetworks



Job title Classification by industry  
(Multi-text Text Classification Task)  
Report

Omneya Essam Kamal

# Overview

You can think of the job industry as the category or general field in which you work. On a job application, "industry" refers to a broad category under which a number of job titles can fall. For example, sales is an industry; job titles under this category can include sales associate, sales manager, manufacturing sales rep, pharmaceutical sales and so on.

## NLP Problem

The problem is a supervised text classification problem, and our goal is to find the most suitable supervised machine learning model to solve it. Given a job title that comes in, we want to assign it to one of 4 industry categories.

## Dataset


The dataset contains of more than 8500 data samples featuring 2 columns:

1- job title    2- industry

```
[132] print('Length of the dataset : {} data samples'.format(len(df)))  
      print('Dataset columns : {}'.format(df.columns))  
  
Length of the dataset : 8586 data samples  
Dataset columns : Index(['job title', 'industry'], dtype='object')
```

As shown in the below figure, there is class imbalance in the data

```
sns.countplot(data=df, x='industry');  
txt="As we can see there is class imbalance here"  
plt.figtext(0.5, -0.1, txt, wrap=True, horizontalalignment='center', fontsize=12);
```



As we can see there is class imbalance here

## Data cleaning

- 1- Nulls: data contains no nulls
- 2- Duplicates: data has more than 4000 duplicates

```
✓ [133] print('Number of duplicates in dataset = {}'.format(df.duplicated().sum()))  
0s  
Number of duplicates in dataset = 4618
```

We will follow 2 approaches:

**First:** not removing duplicates, as this might be changing in real data not duplicated data

**Second:** removing duplicates, as this may bias the ML model

## Text cleaning

```
df['job_title_proc'] = df['job title'].str.lower()  
df['job_title_proc'] = df['job_title_proc'].str.replace(r"(\.|\/)", "")  
df['job_title_proc'] = df['job_title_proc'].str.replace(r'\b-\b', " ")  
df['job_title_proc'] = df['job_title_proc'].str.replace(r'/', " ")  
  
df['job_title_proc'] = df['job_title_proc'].str.split('-').str[0]  
df['job_title_proc'] = df['job_title_proc'].str.split(',').str[0]  
df['job_title_proc'] = df['job_title_proc'].str.replace(r'"', "")  
df['job_title_proc'] = df['job_title_proc'].str.split('-').str[0]  
df['job_title_proc'] = df['job_title_proc'].apply(lambda x: x.replace(']', '').replace('[', ''))  
df['job_title_proc'] = df['job_title_proc'].str.replace('\d+', '')  
df['job_title_proc'] = df['job_title_proc'].str.replace('£', '')  
  
df['job_title_proc'] = df['job_title_proc'].str.replace('{}'.format(string.punctuation), '')  
df['job_title_proc'] = df['job_title_proc'].apply(lambda x: ' '.join([word for word in x.split() if word not in (stopwords)]))  
df['job_title_proc'] = df['job_title_proc'].str.replace('[^a-zA-Z] ', '')
```

- 1- Removing stop words
- 2- Lowering words
- 3- Removing punctuation
- 4- Removing strings between parentheses
- 5- Removing special character and digits
- 6- Splitting composite job title like the example below

|      |                                                            |           |                  |
|------|------------------------------------------------------------|-----------|------------------|
| 8017 | homeroom teacher (native speaker of english)               | Education | homeroom teacher |
| 7822 | senior teacher –young learners- cairo, egypt (egy-s-00034) | Education | senior teacher   |

- 7- Lemmatization

# Models

## Traditional Machine Learning

### 1- Without any data handling

Just using **CountVectorizer**

Testing with different models:

✓  
3s

```
accuracy = []
names = []
for name,model in models:
    model.fit(X_train_dtm, y_train)
    y_pred_class = model.predict(X_test_dtm)
    names.append(name)
    accuracy.append(accuracy_score(y_test, y_pred_class))

for i in range(len(names)):
    print("{} accuracy = {:.3f}".format(names[i],accuracy[i]))
```

```
LogisticRegression accuracy = 0.930
MultinomialNB accuracy = 0.914
SVC accuracy = 0.937
LinearSVC accuracy = 0.932
RandomForest accuracy = 0.923
```

**LinearSVC** is the one I used in deploying Flask API, because it has the highest accuracy

Although I prefer to use **MultinomialNB** as it's computationally very efficient and easy to implement with text data. Plus, it also fits well with multi-class text classification tasks

### 2- Handling imbalanced data

- 1- **SMOTE()**: In a classic oversampling technique, the minority data is duplicated from the minority data population. ... SMOTE works by utilizing a k-nearest neighbour algorithm to create synthetic data.
- 2- **RandomOverSampler()**: over-sample the minority class(es) by picking samples at random with replacement.
- 3- **Class\_weight = "balanced"** : is an argument given in the instance of ML model

```

model = make_pipeline_imb(TfidfVectorizer(), SMOTE(), LogisticRegression(class_weight="balanced"))
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy: ", accuracy_score(y_test, y_pred))
print("Confusion matrix: ", confusion_matrix(y_test, y_pred, labels=labels))

```

```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_i
warnings.warn(msg, category=FutureWarning)
Accuracy: 0.9333954354913834
Confusion matrix: [[1146  27   2  11]
 [ 28 469  21  11]
 [ 11  14 299  11]
 [   2   4   1  90]]

```

Logistic Regression and LinearSVC have the best results of about 93.5

////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////

In my opinion: **F1 score** is the best metric suited to this case, followed by accuracy or AUC-ROC, as we do care about both recall and precision and we are not sensitive to any of them

////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////

### 3- Removing duplicates

Accuracy has decreased after removing duplicates from model, this is expected as the data size has decreased by its half

```

[41] model = make_pipeline_imb(TfidfVectorizer(), SMOTE(), LinearSVC(class_weight='balanced'))
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy: ", accuracy_score(y_test, y_pred))
print("Confusion matrix: ", confusion_matrix(y_test, y_pred, labels=labels))

```

```

Accuracy: 0.8729838709677419
Confusion matrix: [[351  37   5   6]
 [ 25 258   4   7]
 [   9  21 195   5]
 [   1   4   2  62]]

```

```

model = make_pipeline_imb(TfidfVectorizer(), SMOTE(), MultinomialNB())
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy: ", accuracy_score(y_test, y_pred))
print("Confusion matrix: ", confusion_matrix(y_test, y_pred, labels=labels))


plot_confusion_matrix(model, X_test, y_test, values_format='d')
plt.show()

```

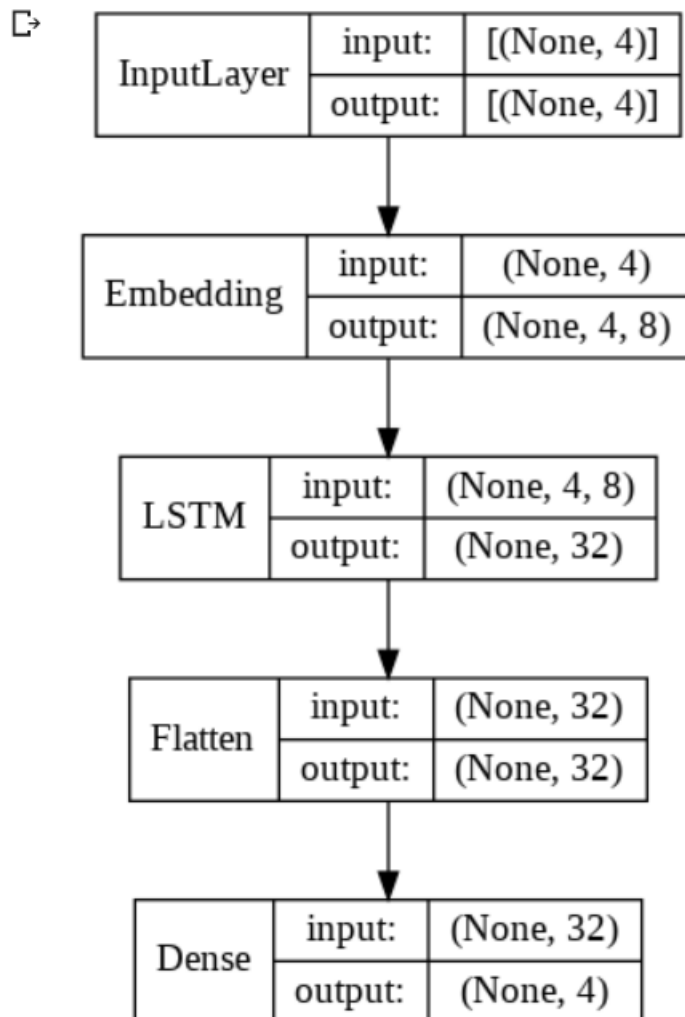
```

Accuracy: 0.8548387096774194
Confusion matrix: [[346  38   3  12]
 [ 21 251   4  18]
 [  9  25 187   9]
 [  1   4   0  64]]

```



## RNN model



- **RNN model is overfitting the data, as the data is very simple and small**
- **How can to extend the model to have better performance?**  
RNN can be restructured to have better performance  
ML models can be tuned with Grid search
- **Limitations of the model:** dataset is very small and the text is too simple as every text sample consists of 2 or 3 words at most

## Flask API

- **Request:** "<http://127.0.0.1:5000/teacher>" (Job title added to request)
- **Response:** "Education" (Predicted Industry by model)

